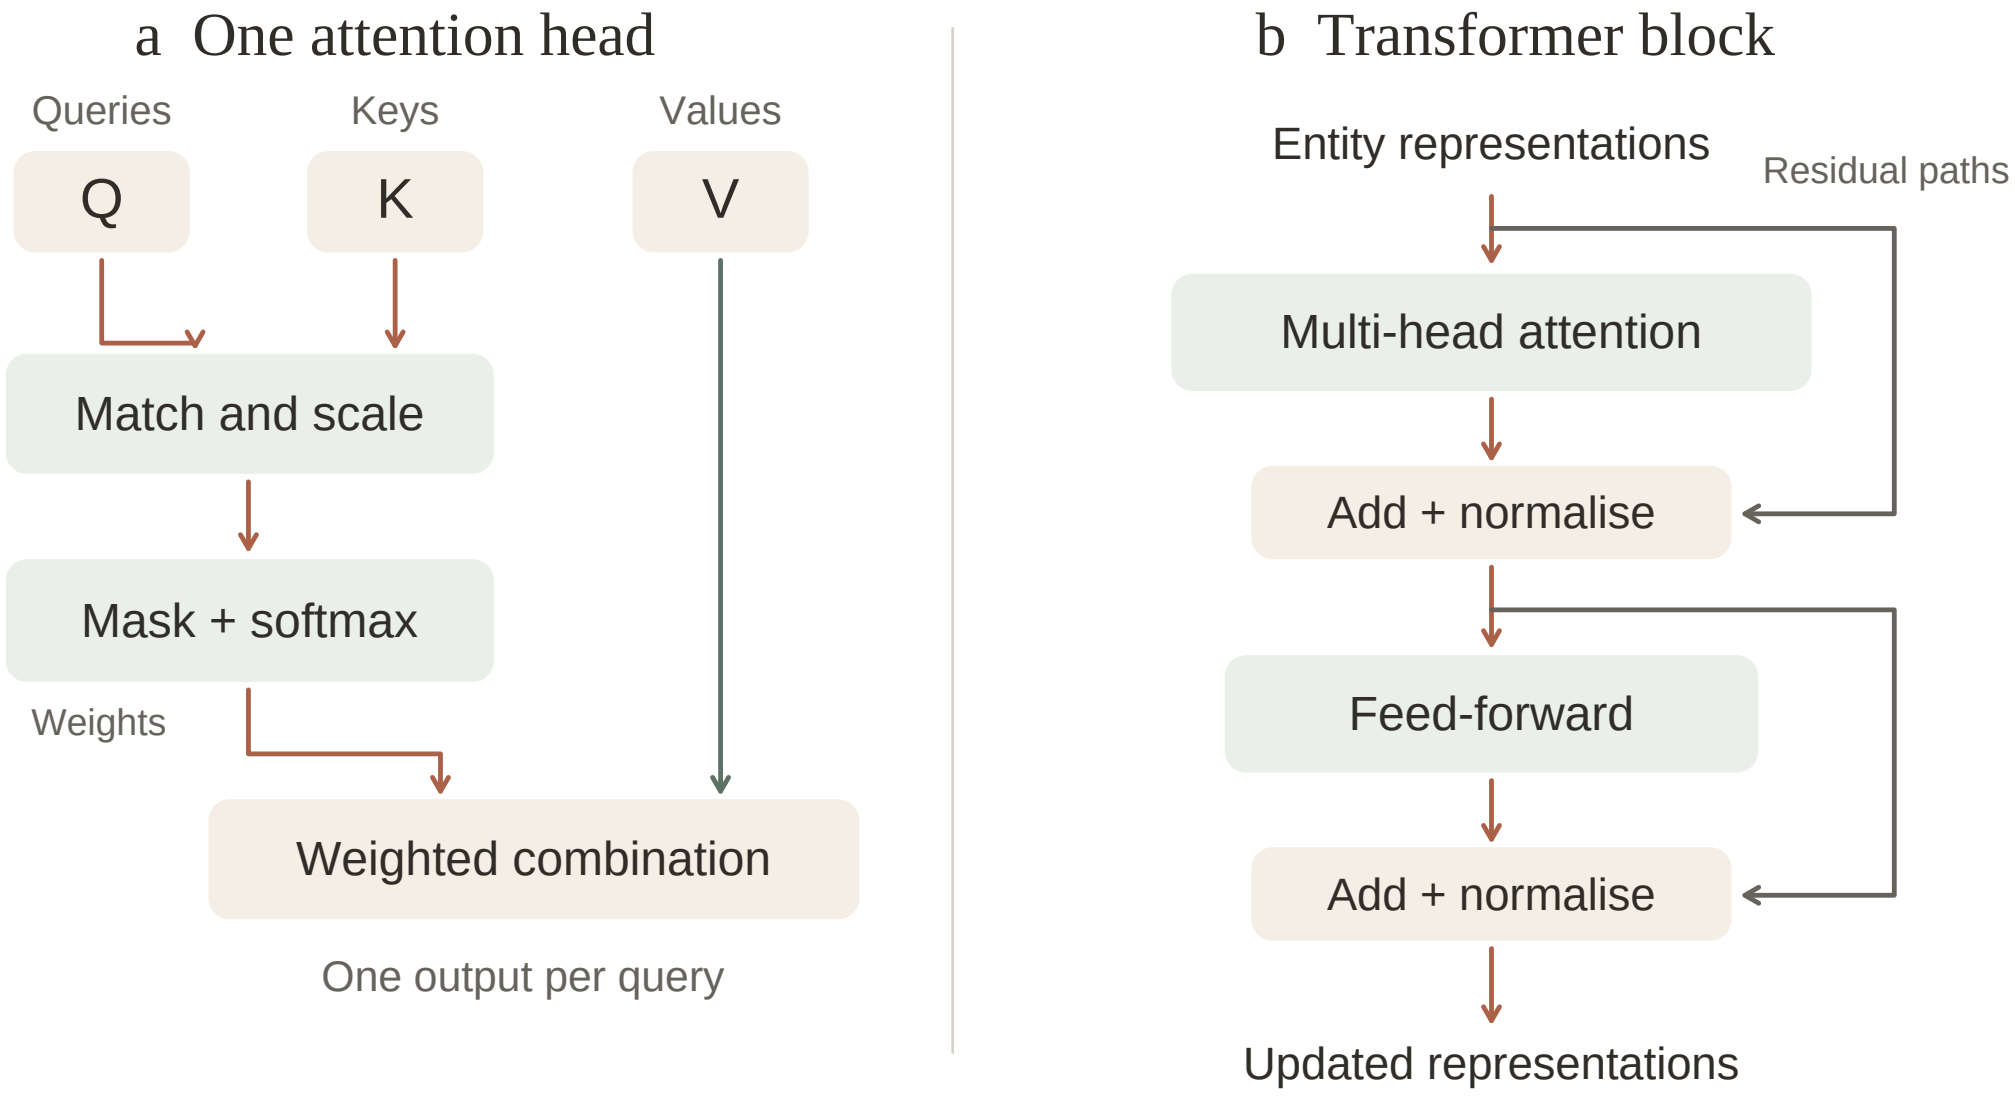
a One attention head
Queries
Keys
Values
Q
K
V
Match and scale
Mask + softmax
Weights
Weighted combination
One output per query
b Transformer block
Entity representations
Residual paths
Multi-head attention
Add + normalise
Feed-forward
Add + normalise
Updated representations