Training: information and learning signals

MADDPG

Replay buffer

Joint observations + actions

Q

Action-value critic

Deterministic actor

Policy gradient through Q

MATD3

Replay buffer

Joint observations + actions

Q1 Q2

Minimum target trains both critics

Deterministic actor

Delayed update through Q1

MAPPO

Recent trajectories

Global / joint information

V

State-value critic

Stochastic actor

Advantages + clipped PPO

Local execution in each agent

Local observation → Trained actor → **Action**